

# Attention is all you need (Transformer)

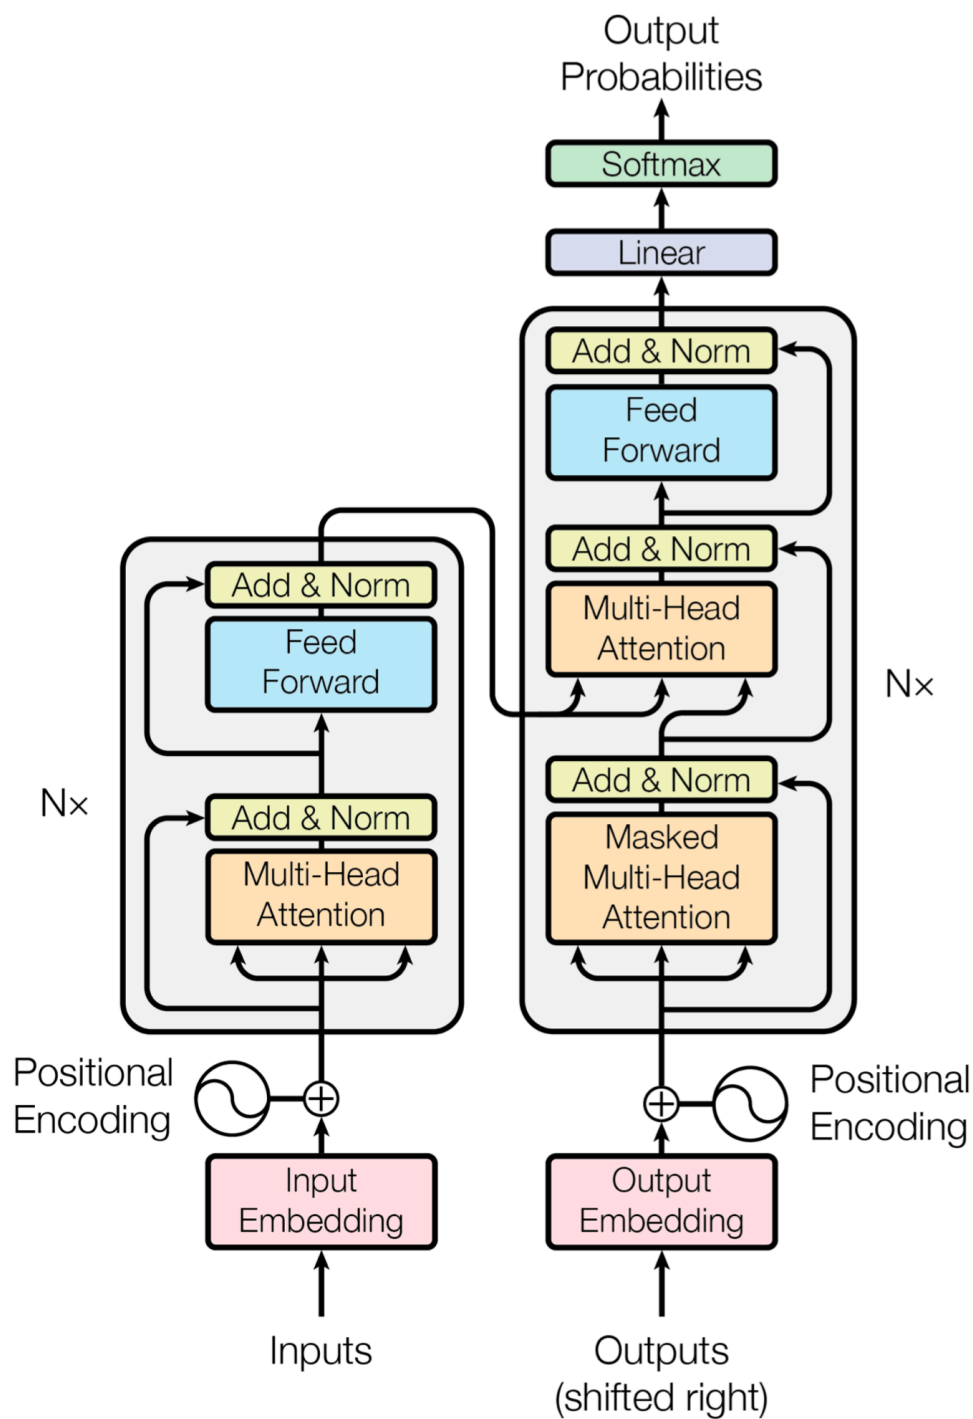


Figure 1: The Transformer - model architecture.

# Attention

An attention function can be described as mapping a query and a set of key-value pairs to an output, where the query, keys, values, and output are all vectors. The output is computed as a weighted sum of the values, where the weight assigned to each value is computed by a compatibility function of the query with the corresponding key.

- 对于某个时刻输出  $y$ ，它在输入  $x$  上各个部分的注意力。可以理解为**权重**。
- 不同机制下的 **attention** 计算方法

Name	Alignment score function	Citation
Additive(*)	$\text{score}(s_t, h_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[s_t; h_i])$	<a href="#">Bahdanau2015</a>
Location-Base	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a s_t)$ Note: This simplifies the softmax alignment max to only depend on the target position.	<a href="#">Luong2015</a>
General	$\text{score}(s_t, h_i) = s_t^\top \mathbf{W}_a h_i$ where $\mathbf{W}_a$ is a trainable weight matrix in the attention layer.	<a href="#">Luong2015</a>
Dot-Product	$\text{score}(s_t, h_i) = s_t^\top h_i$	<a href="#">Luong2015</a>
Scaled Dot-Product(^)	$\text{score}(s_t, h_i) = \frac{s_t^\top h_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where $n$ is the dimension of the source hidden state.	<a href="#">Vaswani2017</a>
Self-Attention(&)	Relating different positions of the same input sequence. Theoretically the self-attention can adopt any score functions above, but just replace the target sequence with the same input sequence.	<a href="#">Cheng2016</a>
Global/Soft	Attending to the entire input state space.	<a href="#">Xu2015</a>
Local/Hard	Attending to the part of input state space; i.e. a patch of the input image.	<a href="#">Xu2015</a> ; <a href="#">Luong2015</a>

- 其中， $s_t$  指的是输出序列的隐藏状态， $h_i$  为输入序列的隐藏状态

## Self-Attention

- 输出序列 为 输入序列

## Transformer中的Multi-head Self-Attention (Dot-product)

$$Attention(Q, K, V) = softmax(\frac{QK^T}{\sqrt{d_k}})V$$

其中,  $d_k$  表示k的维度, paper里默认 64

当k很大时, 得到的点积结果很大, 使得结果处softmax梯度很小, 不利于bp。

- 在encoder的self-attention中, Q, K, V是上一层的encoder输出。对于第一层, 它们是word-embedding和position-embedding相加得到的输出。
- 在decoder的self-attention中, Q、K、V都来自于同一个地方(相等), 它们是上一层decoder的输出。对于第一层decoder, 它们就是word embedding和positional encoding相加得到的输入。但是对于decoder, 我们不希望它能获得下一个time step (即将来的信息), 因此我们需要进行**sequence masking**。
- 在encoder-decoder attention中, Q来自于decoder的上一层的输出, K和V来自于encoder的输出, K和V是一样的。
- Q、K、V三者的维度一样, 即  $d_q = d_k = d_v$

```
import torch
import torch.nn as nn

class ScaledDotProductAttention(nn.Module):
    """Scaled dot-product attention mechanism."""
    def __init__(self, attention_dropout=0.0):
        super(ScaledDotProductAttention,
self).__init__()
        self.dropout = nn.Dropout(attention_dropout)
        self.softmax = nn.Softmax(dim=2)
```

```

def forward(self, q, k, v, scale=None,
attn_mask=None):
    """前向传播.


    Args:
        q: Queries张量, 形状为[B, L_q, D_q]
        k: Keys张量, 形状为[B, L_k, D_k]
        v: Values张量, 形状为[B, L_v, D_v], 一般来说就是
k

        scale: 缩放因子, 一个浮点标量
        attn_mask: Masking张量, 形状为[B, L_q, L_k]


    Returns:
        上下文张量和attention张量
    """
    attention = torch.bmm(q, k.transpose(1, 2))
    if scale:
        attention = attention * scale
    if attn_mask:
        # 给需要mask的地方设置一个负无穷
        attention =
attention.masked_fill_(attn_mask, -np.inf)
    # 计算softmax
    attention = self.softmax(attention)
    # 添加dropout
    attention = self.dropout(attention)
    # 和v做点积
    context = torch.bmm(attention, v)
    return context, attention

```

Scaled Dot-Product Attention



Multi-Head Attention



- 输入的时候Q, K, V在维度上切分, 默认8, 输入512。所以 h 为64。输出的时候再将结果concat, 实验结果比不切割直接通过要优。

## Layer Normalization

- 作用于同一个样本, 计算每一个样本上的均值和方差。

## Feed forward Network

$$FNN(x) = W^T(\max(0, W^T x + b)) + b$$